

In [2]: *### Niraj Neupane*

```
# =====#
# Project 5: Machine Learning for Trading Signal Evaluation (FIXED)
# =====#
# - Downloads SPY data (robust to yfinance column formats)
# - Builds features (no lookahead)
# - Walk-forward ML evaluation (rolling train, forward test)
# - Converts probabilities to a simple long/flat strategy
# - Saves CSV outputs and plots equity curves
#
# Install:
# pip install numpy pandas matplotlib yfinance scikit-learn
# =====#

import os
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt

try:
    import yfinance as yf
except ModuleNotFoundError as e:
    raise ModuleNotFoundError("Missing yfinance. Install with: pip install yfinance")

try:
    from sklearn.pipeline import Pipeline
    from sklearn.preprocessing import StandardScaler
    from sklearn.linear_model import LogisticRegression
    from sklearn.ensemble import RandomForestClassifier
    from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score
except ModuleNotFoundError as e:
    raise ModuleNotFoundError("Missing scikit-learn. Install with: pip install scikit-learn")

plt.style.use("seaborn-v0_8")

# -----
# Config
# -----
OUTPUT_DIR = "outputs"
os.makedirs(OUTPUT_DIR, exist_ok=True)

TICKER = "SPY"
START = "2005-01-01"
END = None

TRAIN_YEARS = 5
TEST_MONTHS = 3
MIN_TRAIN_SAMPLES = 500

PROB_THRESHOLD = 0.52
TCOST_BPS = 2.0
RANDOM_STATE = 42
```

```

# -----
# Robust download helper
# -----
def download_close_series(ticker: str, start: str, end=None) -> pd.Series:
    """
    Robustly return a 1-D Series of Close prices for a single ticker.
    Handles yfinance DataFrame, Series, and MultiIndex columns.
    """
    df = yf.download(ticker, start=start, end=end, auto_adjust=True, progress=False)

    if df is None or len(df) == 0:
        raise ValueError("No data returned from yfinance. Check ticker or internet")

    # Case 1: typical DataFrame with 'Close'
    if isinstance(df, pd.DataFrame):
        # MultiIndex columns sometimes appear, e.g. ('Close', 'SPY')
        if isinstance(df.columns, pd.MultiIndex):
            # Find the 'Close' Level
            if "Close" in df.columns.get_level_values(0):
                close = df["Close"]
                # If still DataFrame (because second Level), select the ticker column
                if isinstance(close, pd.DataFrame):
                    if ticker in close.columns:
                        close = close[ticker]
                    else:
                        # fallback: take first column
                        close = close.iloc[:, 0]
                else:
                    raise ValueError("Downloaded data has MultiIndex columns but no 'Close' column")
            else:
                if "Close" not in df.columns:
                    raise ValueError("Downloaded data does not contain a 'Close' column")
                close = df["Close"]

        # Case 2: Series already
        elif isinstance(df, pd.Series):
            close = df
        else:
            raise ValueError("Unexpected data type returned from yfinance.")

        close = close.dropna()
        close.name = ticker
        return close

# -----
# Feature engineering
# -----
def make_features(close: pd.Series) -> pd.DataFrame:
    """
    Build features using information available at time t to predict y_{t+1}.
    """
    df = pd.DataFrame(index=close.index)
    df["close"] = close

```

```

df["ret_1d"] = df["close"].pct_change()
df["ret_5d"] = df["close"].pct_change(5)
df["ret_21d"] = df["close"].pct_change(21)

df["mom_5d"] = df["ret_1d"].rolling(5).mean()
df["mom_21d"] = df["ret_1d"].rolling(21).mean()
df["mom_63d"] = df["ret_1d"].rolling(63).mean()

df["vol_21d"] = df["ret_1d"].rolling(21).std()
df["vol_63d"] = df["ret_1d"].rolling(63).std()

df["ma_20"] = df["close"].rolling(20).mean()
df["ma_50"] = df["close"].rolling(50).mean()
df["ma_200"] = df["close"].rolling(200).mean()

df["trend_20"] = (df["close"] / df["ma_20"]) - 1
df["trend_50"] = (df["close"] / df["ma_50"]) - 1
df["trend_200"] = (df["close"] / df["ma_200"]) - 1

up = df["ret_1d"].clip(lower=0)
down = (-df["ret_1d"]).clip(lower=0)
rs = up.rolling(14).mean() / (down.rolling(14).mean() + 1e-12)
df["rsi_14"] = 100 - (100 / (1 + rs))

# Target: next-day return and direction
df["fwd_ret_1d"] = df["ret_1d"].shift(-1)
df["y"] = (df["fwd_ret_1d"] > 0).astype(int)

df = df.dropna()
return df

# -----
# Walk-forward splits
# -----
def walk_forward_splits(index: pd.DatetimeIndex, train_years: int, test_months: int):
    start_date = index.min()
    end_date = index.max()
    test_start = start_date + pd.DateOffset(years=train_years)

    while test_start < end_date:
        train_start = test_start - pd.DateOffset(years=train_years)
        train_end = test_start - pd.DateOffset(days=1)

        test_end = test_start + pd.DateOffset(months=test_months) - pd.DateOffset(d)
        yield train_start, train_end, test_start, min(test_end, end_date)

        test_start = test_start + pd.DateOffset(months=test_months)

# -----
# Backtest utilities
# -----
def perf_metrics(returns: pd.Series) -> dict:
    r = returns.dropna()
    if len(r) < 50:

```

```

        return {"n": int(len(r))}

    ann_ret = (1 + r).prod() ** (252 / len(r)) - 1
    ann_vol = r.std() * np.sqrt(252)
    sharpe = np.nan if ann_vol == 0 else ann_ret / ann_vol

    eq = (1 + r).cumprod()
    dd = eq / eq.cummax() - 1
    mdd = dd.min()

    win_rate = (r > 0).mean()

    return {
        "n": int(len(r)),
        "ann_return": float(ann_ret),
        "ann_vol": float(ann_vol),
        "sharpe": float(sharpe) if np.isfinite(sharpe) else np.nan,
        "max_drawdown": float(mdd),
        "win_rate": float(win_rate),
    }

def apply_transaction_costs(position: pd.Series, returns: pd.Series, tcost_bps: float):
    tcost = tcost_bps / 10000.0
    turnover = (position - position.shift(1)).abs().fillna(0.0)
    cost = turnover * tcost
    strat = position.shift(1).fillna(0.0) * returns - cost
    return strat

# -----
# Main
# -----
def main():
    print("Downloading data...")
    close = download_close_series(TICKER, START, END)

    data = make_features(close)
    feature_cols = [
        "ret_1d", "ret_5d", "ret_21d",
        "mom_5d", "mom_21d", "mom_63d",
        "vol_21d", "vol_63d",
        "trend_20", "trend_50", "trend_200",
        "rsi_14",
    ]

    X_all = data[feature_cols].copy()
    y_all = data["y"].copy()
    fwd_ret = data["fwd_ret_1d"].copy()
    idx = data.index

    logit = Pipeline([
        ("scaler", StandardScaler()),
        ("clf", LogisticRegression(max_iter=2000, random_state=RANDOM_STATE))
    ])

```

```

rf = RandomForestClassifier(
    n_estimators=400,
    max_depth=6,
    min_samples_leaf=50,
    random_state=RANDOM_STATE,
    n_jobs=-1
)

rows = []

for train_start, train_end, test_start, test_end in walk_forward_splits(idx, TR
    train_mask = (idx >= train_start) & (idx <= train_end)
    test_mask = (idx >= test_start) & (idx <= test_end)

    X_train, y_train = X_all.loc[train_mask], y_all.loc[train_mask]
    X_test, y_test = X_all.loc[test_mask], y_all.loc[test_mask]

    if len(X_train) < MIN_TRAIN_SAMPLES or len(X_test) < 20:
        continue

    logit.fit(X_train, y_train)
    rf.fit(X_train, y_train)

    p_logit = logit.predict_proba(X_test)[:, 1]
    p_rf = rf.predict_proba(X_test)[:, 1]

    tmp = pd.DataFrame(index=X_test.index)
    tmp["y_true"] = y_test.values
    tmp["fwd_ret"] = fwd_ret.loc[X_test.index].values
    tmp["p_logit"] = p_logit
    tmp["p_rf"] = p_rf
    rows.append(tmp)

pred = pd.concat(rows).sort_index()

if pred.empty:
    raise ValueError("No walk-forward predictions were produced. Try reducing T")

pred["pred_logit"] = (pred["p_logit"] >= 0.5).astype(int)
pred["pred_rf"] = (pred["p_rf"] >= 0.5).astype(int)

pred["pos_logit"] = (pred["p_logit"] >= PROB_THRESHOLD).astype(float)
pred["pos_rf"] = (pred["p_rf"] >= PROB_THRESHOLD).astype(float)

pred["ret_logit"] = apply_transaction_costs(pred["pos_logit"], pred["fwd_ret"])
pred["ret_rf"] = apply_transaction_costs(pred["pos_rf"], pred["fwd_ret"], TCOST)

pred["ret_buy_hold"] = pred["fwd_ret"]

mom_signal = (data["ret_21d"].reindex(pred.index) > 0).astype(float)
pred["ret_momentum"] = apply_transaction_costs(mom_signal, pred["fwd_ret"], TCO)

pred["eq_logit"] = (1 + pred["ret_logit"]).fillna(0)).cumprod()
pred["eq_rf"] = (1 + pred["ret_rf"]).fillna(0)).cumprod()
pred["eq_buy_hold"] = (1 + pred["ret_buy_hold"]).fillna(0)).cumprod()
pred["eq_momentum"] = (1 + pred["ret_momentum"]).fillna(0)).cumprod()

```

```

def cls_metrics(y_true, prob, pred_label):
    out = {
        "accuracy": accuracy_score(y_true, pred_label),
        "precision": precision_score(y_true, pred_label, zero_division=0),
        "recall": recall_score(y_true, pred_label, zero_division=0),
        "f1": f1_score(y_true, pred_label, zero_division=0),
    }
    try:
        out["roc_auc"] = roc_auc_score(y_true, prob)
    except Exception:
        out["roc_auc"] = np.nan
    return out

m_logit = cls_metrics(pred["y_true"], pred["p_logit"], pred["pred_logit"])
m_rf = cls_metrics(pred["y_true"], pred["p_rf"], pred["pred_rf"])

t_logit = perf_metrics(pred["ret_logit"])
t_rf = perf_metrics(pred["ret_rf"])
t_bh = perf_metrics(pred["ret_buy_hold"])
t_mom = perf_metrics(pred["ret_momentum"])

metrics_df = pd.DataFrame([
    {"model": "LogisticRegression", **m_logit, **{f"trade_{k)": v for k, v in t_logit.items()}},
    {"model": "RandomForest", **m_rf, **{f"trade_{k)": v for k, v in t_rf.items()}},
    {"model": "BuyHold", **{f"trade_{k)": v for k, v in t_bh.items()}},
    {"model": "Momentum_21d", **{f"trade_{k)": v for k, v in t_mom.items()}},
])
print("\n===== OUT-OF-SAMPLE METRICS =====")
print(metrics_df)

pred.to_csv(os.path.join(OUTPUT_DIR, "ml_predictions.csv"))
pred[["eq_logit", "eq_rf", "eq_buy_hold", "eq_momentum"]].to_csv(os.path.join(OUTPUT_DIR, "ml_predictions.csv"))
metrics_df.to_csv(os.path.join(OUTPUT_DIR, "ml_metrics.csv"), index=False)

print("\nSaved outputs to ./outputs")

plt.figure(figsize=(11, 6))
pred["eq_buy_hold"].plot(label="Buy & Hold")
pred["eq_momentum"].plot(label="Momentum (21d)")
pred["eq_logit"].plot(label=f"Logit (thr={PROB_THRESHOLD})")
pred["eq_rf"].plot(label=f"RF (thr={PROB_THRESHOLD})")
plt.title("Project 5 – ML Trading Signals (Walk-Forward OOS Equity Curves)")
plt.ylabel("Growth of $1")
plt.legend()
plt.tight_layout()
plt.show()

# In Jupyter, just run main()
# In .py file, it will run automatically
if __name__ == "__main__":
    main()

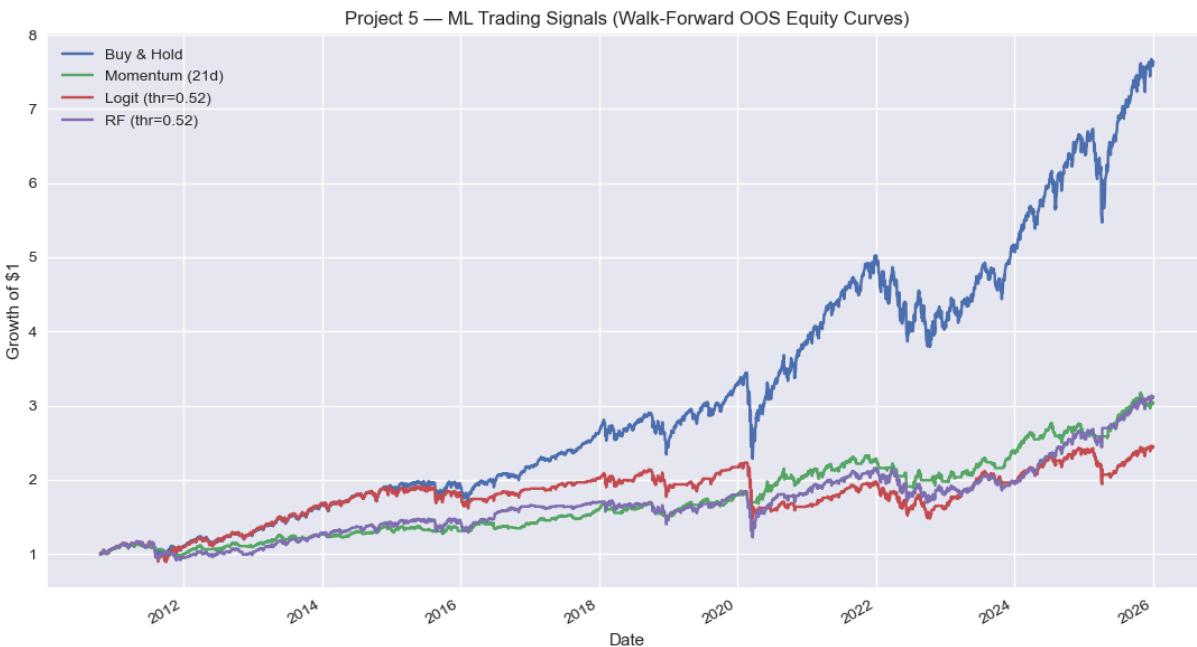
```

Downloading data...

===== OUT-OF-SAMPLE METRICS =====

	model	accuracy	precision	recall	f1	roc_auc	\
0	LogisticRegression	0.538160	0.555486	0.837418	0.66792	0.493419	
1	RandomForest	0.542342	0.556400	0.862394	0.67640	0.490348	
2	BuyHold	NaN	NaN	NaN	NaN	NaN	
3	Momentum_21d	NaN	NaN	NaN	NaN	NaN	
	trade_n	trade_ann_return	trade_ann_vol	trade_sharpe	trade_max_drawdown		\
0	3826	0.060734	0.152375	0.398585	-0.340487		
1	3826	0.077849	0.154815	0.502850	-0.337173		
2	3826	0.143317	0.170992	0.838151	-0.337173		
3	3826	0.076024	0.106298	0.715193	-0.187479		
	trade_win_rate						
0		0.404861					
1		0.405907					
2		0.554626					
3		0.385781					

Saved outputs to ./outputs



In []: